

# Instantiations Must Go!

Matt Calabrese  
Zach Laine

May 10, 2010

A funny thing happened  
on the way to BoostCon

# Why did you tell me this ridiculous story?

- Like Kennybrook, we did something stupid with serendipitous results

# Why did you tell me this ridiculous story?

- This isn't a “serious business” BoostCon lecture

# Why did you tell me this ridiculous story?

- I just like telling it

# The Problem

# The Problem

- Template Metaprogramming is fun
  - Allows you to do lots of things at compile time that make your life easier/better
  - Makes your brain seem bigger to others

# The Problem

- Template instantiations are expensive at compile-time



# The Problem

- Each instantiation must do a lot of work inherently
- Some compilers have really slow instantiation code (I'm looking at you, GCC < 4.5)

# The Problem

- Notable Boost examples:
  - Proto uses macro workarounds extensively
  - Spirit code can be unmaintainable in large enough quantities

The “Solution”

# The Solution

- Easy – don't instantiate templates...
- ... which is stupid, remember?

# The Solution

- Here's what we want to avoid:

```
template <typename T>
struct metafunction
{
    typedef /* something with T */ type;
};
```

```
metafunction<int>::type x = /* ... */;
```

- Just don't reach into `metafunction` to grab `type`, and we don't instantiate any templates, right?
- Not really (more on this later)

# The Solution: The General Idea

- C++0x's `decltype` and trailing return type function declaration syntax make it possible to do a lot of TMP without struct templates – let's try some approaches that use expressions passed to `decltype`.

# The Solution: The General Idea

- `decltype`, in loose terms, yields the type of the expression given to it, turning this:

```
template <typename T, typename U>
struct result_of_plus
{ typedef /* usually something complicated */ type; };
```

```
template <typename T, typename U>
typename result_of_plus<T, U>::type
operator+(const T& t, const U& u)
{ return t + u; }
```

- ... into this:

```
template <typename T, typename U>
decltype((*T*)0 + (*U*)0)
operator+(const T& t, const U& u)
{ return t + u; }
```

# The Solution: The General Idea

- The trailing return type syntax allows us to use parameter names in `decltype` expressions. This:

```
template <typename T, typename U>
decltype((*T*)0) + ((*U*)0)
operator+(const T& t, const U& u)
{ return t + u; }
```

- ... becomes this:

```
template <typename T, typename U>
auto operator+(const T& t, const U& u) -> decltype(t + u)
{ return t + u; }
```



# The Solution: Approach 1

```
// vector.h:

template <
    typename ... Types
>
struct vector_;

template <>
struct vector_<>
{};

template <
    typename Head,
    typename ... Tail
>
struct vector_<Head, Tail ...> :
    vector_<Tail ...>
{};
```

# The Solution: Approach 1

```
#include <boost/mpl/print.hpp>
#include "vector.h"

template <typename ... Args, typename T>
vector_<Args ..., T>
(* push_back_(vector_<Args ...> (*) (), T (*) ())) ();

int main()
{
    typedef vector_<> empty_vec;

    typedef decltype(
        push_back_(
            empty_vec (*) (), (int (*) ()) 0
        ) ()
    ) int_vec;

    typedef boost::mpl::print<int_vec>::type foo;

    return 0;
}
```

## Compiler Output:

```
.../boost/mpl/print.hpp: In instantiation of
boost::mpl::print<vector_<int> >: ...
```

# The Solution: Approach 1

- Pros:
  - It works!
- Cons:
  - It's disgusting!

# The Solution: Approach 2: Wrap/Unwrap

- Let's automate this wrapping business

```
// wrap_unwrap.h
```

```
template <typename T>  
struct wrap;
```

```
template <typename T>  
struct unwrap;
```

```
template <typename T>  
struct unwrap<wrap<T> *>  
{ typedef T type; };
```

# The Solution: Approach 2, Part II

```
#include <boost/mpl/print.hpp>
#include "vector.h"
#include "wrap_unwrap.h"

template <typename ... Args, typename T>
wrap<vector_<Args ..., T> > *
push_back_(wrap<vector_<Args ...> > *, wrap<T> *);

int main()
{
    typedef vector_<> empty_vec;

    typedef unwrap<
        decltype(
            push_back_(
                (wrap<empty_vec> *)0,
                (wrap<int> *)0
            )
        )
    >::type int_vec;

    typedef boost::mpl::print<int_vec>::type foo;

    return 0;
}
```

## Compiler Output:

```
.../boost/mpl/print.hpp: In instantiation of boost::mpl::print<vector_<int>
>: ...
```

# The Solution: Approach 2

- Pros:
  - It still works!
- Cons:
  - It's slightly less disgusting!

# The Solution: The final version's wrappers

```
// wrap_unwrap_2.h

template< typename T >
struct type_ {};

template< typename T >
type_< T >& type();

template< typename T >
struct unwrap;

template< typename T >
struct unwrap< type_<T>& >
{ typedef T type; };
```

- We've established a new parameter/return type convention, that:
  - Leaves types incomplete where possible
  - Requires `type()` and `unwrap()` calls, but only at the library boundaries

# The Solution: The final version

```
#include <boost/mpl/print.hpp>
#include "vector.h"
#include "wrap_unwrap_2.h"

template< typename ... T >
type_< vector_< T... > >& vector();

template< typename ... T, typename Elem >
type_< vector_< T..., type_< Elem >& > >&
push_back( type_< vector_< T... > >&, type_< Elem >& );

int main()
{
    typedef unwrap<
        decltype(
            push_back(vector(), type<int>())
        )
    >::type int_vec;

    typedef boost::mpl::print<int_vec>::type foo;

    return 0;
}
```

## Compiler Output:

```
.../boost/mpl/print.hpp: In instantiation of
boost::mpl::print<vector_<type_<int>& > >: ...
```



# The Solution: The final version

- Pros:
  - As designed, it does TMP without instantiating `struct` definitions
  - It's more readable than standard TMP
    - Allows mere mortals to read TMP code written by others
    - Is much more teachable
  - It's damn sexy

# The Solution: The final version

- Cons:
  - It still instantiates `struct` declarations
  - Its compile-time-reducing performance is questionable

# The Solution: The numbers

- Compile-time performance measures:
  - Test was a deeply nested series of vector `push_back`'s
  - A robust test would require reimplementing Boost.MPL and Spirit, Proto, or some other MPL user

# The Solution: The numbers

- GCC  $\geq$  4.4 Numbers

gcc (64-bit) 4.4.0:	mpl: 0.5575s	ftmpl: 0.4467s	-- 1.25x	ftmpl speedup
gcc (64-bit) 4.4.1:	mpl: 0.5726s	ftmpl: 0.4419s	-- 1.30x	ftmpl speedup
gcc (64-bit) 4.4.2:	mpl: 0.5628s	ftmpl: 0.4409s	-- 1.28x	ftmpl speedup
gcc (32-bit) 4.4.3:	mpl: 0.5777s	ftmpl: 0.4559s	-- 1.27x	ftmpl speedup
gcc (64-bit) 4.4.4:	mpl: 0.5679s	ftmpl: 0.4592s	-- 1.24x	ftmpl speedup
gcc (64-bit) 4.5.0:	mpl: 0.5116s	ftmpl: 0.4773s	-- 1.07x	ftmpl speedup

# FTMPL – “Function Template MPL”

- The FTMPL version:

```
#include <boost/mpl/print.hpp>

#include "vector.hpp"
#include "type.hpp"
#include "unwrap.hpp"

int main()
{
    using namespace boost::ftmpl;

    typedef BOOST_FTMPL_UNWRAP (
        push_back(vector(), type<int>())
    ) int_vec;

    typedef boost::mpl::print<int_vec>::type foo;

    return 0;
}
```

# FTMPL: Readability

- `decltype` allows us to do some very interesting TMP with expressions:

```
typedef BOOST_FTMPL_UNWRAP(  
    int_<3>() + int_<6>()  
) int_9;  
typedef boost::mpl::print<int_9>::type foo;  
  
int result = BOOST_FTMPL_UNWRAP_VALUE(  
    int_<3>() & int_<6>()  
);  
std::cout << "result = " << result << '\n';
```

## Compiler Output:

```
.../boost/mpl/print.hpp: In instantiation of  
boost::mpl::print<boost::ftmpl::value_t<int, 9> >: ...
```

## Program Output:

```
result = 2
```

- Metafunction objects actually look like ...  
function objects!

```
struct is_same_t
{
    template< typename L, typename R >
    type_< false_t >& operator()( L&, R& );

    template< typename L >
    type_< true_t >& operator()( L&, L& );
};

typedef BOOST_FTMPL_UNWRAP (
    is_same_t()(type< char >(), type< char >())
) is_same_result;
typedef boost::mpl::print<is_same_result>::type foo;
```

## Compiler Output:

```
.../boost/mpl/print.hpp: In instantiation of
boost::mpl::print<boost::ftmpl::value_t<bool, true> >: ...
```

- Contrast this with `type_traits/is_same.hpp`:

```
template <typename T>
::boost::type_traits::yes_type
BOOST_TT_DECL is_same_tester(T*, T*);

::boost::type_traits::no_type
BOOST_TT_DECL is_same_tester(...);

template <typename T, typename U>
struct is_same_impl
{
    static T t;
    static U u;

    BOOST_STATIC_CONSTANT(bool, value =
        (::boost::type_traits::ice_and<
            (sizeof(type_traits::yes_type) == sizeof(detail::is_same_tester(&t, &u))),
            (::boost::is_reference<T>::value == ::boost::is_reference<U>::value),
            (sizeof(T) == sizeof(U))
        >::value));
};
```



- Higher-order metafunctions are actually simple enough to grok at first glance:

```
template< typename Metafun, typename ... Args >
auto apply( type_< Metafun >&, Args& ... args )
    -> decltype( Metafun()( args... ) );

extern type_< is_same_t >& is_same;

typedef BOOST_FTMPL_UNWRAP(
    apply( is_same, type< char >(), type< int >() )
) applied_is_same_result;
typedef boost::mpl::print<applied_is_same_result>::type foo;
```

## Compiler Output:

```
.../boost/mpl/print.hpp: In instantiation of
boost::mpl::print<boost::ftmpl::value_t<bool, false> >: ...
```

# FTMPL: Readability

- Higher-order metafunctions are actually simple enough to grok at first glance:

```
template<
    typename Fn,
    typename State
>
State& fold( Fn&, State&, type_< vector_t<> >& );

template<
    typename Fn,
    typename State,
    typename Head,
    typename ... Tail
>
auto fold(
    type_< Fn >&,
    type_< State >&,
    type_< vector_t< type_< Head >, Tail... > >&
)
-> decltype(
    fold(
        type_< Fn >(),
        Fn()( type_< State >(), type_< Head >() ),
        vector_< Tail... >()
    )
);
```

- Contrast with mpl/fold.hpp (main implementation is elsewhere):

```
template<
    typename BOOST_MPL_AUX_NA_PARAM(Sequence)
    , typename BOOST_MPL_AUX_NA_PARAM(State)
    , typename BOOST_MPL_AUX_NA_PARAM(ForwardOp)
>
struct fold
{
    typedef typename aux::fold_impl<
        ::boost::mpl::O1_size<Sequence>::value
        , typename begin<Sequence>::type
        , typename end<Sequence>::type
        , State
        , ForwardOp
        >::state type;

    BOOST_MPL_AUX_LAMBDA_SUPPORT(3, fold, (Sequence, State, ForwardOp))
};
```

# Functional MPL Work-In-Progress

- FTMPL (quite broken)
  - <https://svn.boost.org/svn/boost/sandbox/ftmpl>
- MPL 0x (maybe broken?)
  - <http://github.com/dabrahams/mpl0x>
  - <http://github.com/swatanabe/mpl0x>

Questions?