



BoostCon 2010

# Multithreaded C++0x: the Dawn of a new Standard

Michael Wong  
michaelw@ca.ibm.com  
IBM Toronto Lab  
Canadian C++ Standard Committee

# IBM Rational Disclaimer

- **© Copyright IBM Corporation 2010. All rights reserved. The information contained in these materials is provided for informational purposes only, and is provided AS IS without warranty of any kind, express or implied. IBM shall not be responsible for any damages arising out of the use of, or otherwise related to, these materials. Nothing contained in these materials is intended to, nor shall have the effect of, creating any warranties or representations from IBM or its suppliers or licensors, or altering the terms and conditions of the applicable license agreement governing the use of IBM software. References in these materials to IBM products, programs, or services do not imply that they will be available in all countries in which IBM operates. Product release dates and/or capabilities referenced in these materials may change at any time at IBM's sole discretion based on market opportunities or other factors, and are not intended to be a commitment to future product or feature availability in any way. IBM, the IBM logo, Rational, the Rational logo, Telelogic, the Telelogic logo, and other IBM products and services are trademarks of the International Business Machines Corporation, in the United States, other countries or both. Other company, product, or service names may be trademarks or service marks of others.**

# The IBM Rational C/C++ Café

[ibm.com/rational/cafe/community/ccpp](http://ibm.com/rational/cafe/community/ccpp)

The screenshot shows the IBM Rational C/C++ Café website. At the top, there is a navigation bar with the IBM logo, a welcome message for a guest, a link to sign in or register, and a search bar. Below this is a large green banner with the text "C/C++ Café" and "Boosting Performance, Productivity, and Portability". Under the banner is a horizontal menu with tabs for Cafés, Resource Library, Discussion Forums, Blogs, Products, Standards, and Platform Partners. The "Cafés" tab is selected, showing a list of links: Articles, Presentations, Documentation, Downloads, Learn, and Support. The "Discussion Forums" tab shows links for C/C++ General, C/C++ Language Star, and Cafe Feedback. The "Blogs" tab shows a list of blog titles: "The C/C++ Market Place: Product Management", "Commercial Computing with C/C++", "Parallel and Multi-Core Computing with C/C++", "Scientific Computing with C/C++", "C Standard", and "C++ Standard". A "Welcome to the C/C++ Café" message is visible at the bottom of the main content area.

IBM. Welcome, Guest | [Sign in or register](#)

## C/C++ Café

Boosting Performance, Productivity, and Portability

Cafés	Resource Library	Discussion Forums	Blogs	Products	Standards	Platform Partners
<a href="#">IBM Rational C/C++ Overview</a>	<a href="#">Articles, Presentations</a>	<a href="#">C/C++ General</a>	<a href="#">The C/C++ Market Place: Product Management</a>			
	<a href="#">Documentation</a>	<a href="#">C/C++ Language Star</a>	<a href="#">Commercial Computing with C/C++</a>			
	<a href="#">Downloads</a>	<a href="#">Cafe Feedback</a>	<a href="#">Parallel and Multi-Core Computing with C/C++</a>			
	<a href="#">Learn</a>		<a href="#">Scientific Computing with C/C++</a>			
	<a href="#">Support</a>		<a href="#">C Standard</a>			
			<a href="#">C++ Standard</a>			

Welcome to the C/C++ Café



### Join

A community of Industry Leaders in C/C++ Technology



### Download

Trials of new technology



### Learn

To Take full advantage of the IBM C/C++ compilers



### Share

Participate in forum discussions. Follow and Respond to Blogs.

# Agenda

- **Concurrent C++0x examples**
- Atomics header
- Atomic types
- Atomic operations
- Atomic relations
- Atomic ordering
- Q/A

# Hello Concurrent World

```
#include <iostream>
#include <thread> //#1
void hello() //#2
{
    std::cout<<"Hello Concurrent World"<<std::endl;
}
int main()
{
    std::thread t(hello); //#3
    t.join(); //#4
}
```

# Is this valid C++ today? Are these equivalent?

```
int x = 0;
atomic<int> y = 0;

Thread 1:
  x = 17;
  y.store(1,
memory_order_release);
  // or:      y.store(1);

Thread 2:
  while
  (y.load(memory_order_acquire) != 1)
  // or:      while
  (y.load() != 1)

  assert(x == 17);
```

```
int x = 0;
atomic<int> y = 0;

Thread 1:
  x = 17;
  y = 1;

Thread 2:
  while (y != 1)
    continue;
  assert(x == 17);
```

# Memory Model

- **Locks and atomic operations communicate non-atomic writes between two threads**
- **Volatile is not atomics**
- **Memory races cause undefined behavior**
- **Some optimizations are no longer legal**
- **Compiler may assume some loops terminate**



# Message shared memory

- **Writes are explicitly communicated**
  - Between pairs of threads
  - Through a lock or an atomic variable
- **The mechanism is acquire and release**
  - One thread releases its memory writes
    - `V=32; atomic_store_explicit(&a,3, memory_order_release );`
  - Another thread acquires those writes
    - `i=atomic_load_explicit(&a, memory_order_acquire ); i+v;`



# What is a memory location

- **A non-bitfield primitive data object**
- **A sequence of adjacent bitfields**
  - Not separated by a structure boundary
  - Not interrupted by the null bitfield
  - Avoid expensive atomic read-modify-write operations on bitfields

# Data race condition

- **A non-atomic write to a memory location in one thread**
- **A non-atomic read from or write to that same location in another thread**
- **With no happens-before relations between them**
- **Is undefined behaviour**

# Effect on compiler optimization

- **Some rare optimizations are restricted**
  - Fewer speculative writes
  - Fewer speculative reads
- **Some common optimizations can be augmented**
  - They may assume that loops terminate
  - Nearly always true

# Atomics: To Volatile or Not Volatile

- **Too much history in volatile to change its meaning**
- **It is not used to indicate atomicity like Java**
- **Volatile atomic means something from the environment may also change this in addition to another thread**

# Requirements on atomics

- **Static initialization**
- **Reasonable implementation on current hardware**
- **Relative novices can write working code**
- **Experts can performance efficient code**

# Consistency problem

- **X and y are atomic and initially 0**
  - Thread 1: `x=1;`
  - Thread 2: `y=1;`
  - Thread 3: `if (x==1 && y==0)`
  - Thread 4: `if ( x==0 && y==1)`
- **Are both conditions exclusive?**
  - Is there a total store order?
- **The hardware/software system may not provide it**
- **Programming is harder without it**

# Consistency models

- **Sequentially consistent**

- What is observed is consistent with a sequential ordering of all events in the system
  - But comes with a very heavy cost

- **Weaker models**

- More complex to code for some
  - But very efficient

- **What we decided**

- Default is sequential consistency
- But allow weaker semantics explicitly



# Atomic Library (N2427)

- The problem:

- **Would like to implement, for example, counters, without locks using atomics**

```
atomic<int> x;  
void increment() {  
    ++x;  
}
```

```
int x;  
mutex m;  
void increment(){
```

```
    lock_guard_(m);  
    x=x+1;  
}
```

- **Advantages**

- Sometimes enables much better performance
- No space for locks, sometimes simpler.
- Potentially safe for use with signal handlers, across processes.

# Atomic DCL

```
T x;  
atomic_bool x_init(false);  
mutex m;  
if (!x_init) {  
    lock_guard _(m);  
    if (!x_init) {  
        x = ....  
        x_init = true;  
    }  
}  
use x;
```

- Note: Atomics are still tricky. Only a single memory operation at a time is atomic!

# Atomic Design

- Want shared variables
  - **that can be concurrently updated without introducing data race,**
  - **that are atomically updated and read**
    - **half updated states are not visible,**
- that are implemented without lock overhead whenever the hardware allows,
- that provide access to hardware atomic read-modify write (fetch-and-add, xchg, cmpxchg) instructions whenever possible.

## Race Free semantics and Atomic Memory operations

- **If a program has a race, it has undefined behavior**
  - This is sometimes known as “catch fire” semantics
  - No compiler transformation is allowed to introduce a race
    - no invented writes
    - Possibly fewer speculative stores and **(potentially) loads**
- **There are atomic memory operations that don't cause races**
  - Can be used to implement locks/mutexes
  - Also useful for lock-free algorithms
- **Atomic memory operations are expressed as library function calls**
  - Reduces need for new language syntax

# Atomic Operations and Type

- **Data race: if there is no enforced ordering between two accesses to a single memory location from separate threads, one or both of those accesses is not atomic, and one or both is a write, this is a data race, and causes undefined behavior.**
- **These types avoid undefined behavior and provide an ordering of operations between threads**

# Standard Atomic Types

- **#include <stdatomic>**
- **atomic\_flag**
- **atomic\_bool**
- **atomic\_address**
- **Integral types:**
  - atomic\_char, atomic\_schar, atomic\_uchar, atomic\_short, atomic\_ushort, atomic\_int, atomic\_uint, atomic\_long, atomic\_ulong, atomic\_llong, atomic\_ullong, atomic\_char16\_t, atomic\_char32\_t, atomic\_wchar\_t
- **Typedefs like those in <stdint>**
  - atomic\_int\_least8\_t, atomic\_uint\_least8\_t, atomic\_int\_least16\_t, atomic\_uint\_least16\_t, atomic\_int\_least32\_t, atomic\_uint\_least32\_t, atomic\_int\_least64\_t, atomic\_uint\_least64\_t, atomic\_int\_fast8\_t, atomic\_uint\_fast8\_t, atomic\_int\_fast16\_t, atomic\_uint\_fast16\_t, atomic\_int\_fast32\_t, atomic\_uint\_fast32\_t, atomic\_int\_fast64\_t, atomic\_uint\_fast64\_t, atomic\_intptr\_t, atomic\_uintptr\_t, atomic\_size\_t, atomic\_ssize\_t, atomic\_ptrdiff\_t, atomic\_intmax\_t, atomic\_uintmax\_t
- **is\_lock\_free();**
- **Non-copyable, non-assignable**

# Minimal atomics

- **Need 1 primitive data types that is a must, most modern hardware has instructions to implement the atomic operations**
  - **for small types**
  - **and bit-wise comparison, assignment (which we require)**
  - atomic\_flag type

```
static std::atomic_flag v1= ATOMIC_FLAG_INIT
If (atomic_flag_test_and_set(&v1))
    atomic_flag_clear(&v1);
```
- **For other types, hardware, atomic operations may be emulated with locks.**
  - Sometimes this isn't good enough:
    - across processes, in signal/interrupt handlers.
  - is\_lock\_free() returns false if locks are used, and operations may block.
- **Operations on variable have attributes, which can be explicit**
  - Acquire=get other memory writes
  - Release=give my memory writes
  - Acq\_and\_rel=Acquire and release at the same time
  - Relaxed=no acquire or released, non-deterministic, not synchronizing with the rest of memory, but still sequential view of that variable
  - Seq-cst=Fully ordered,extra ordering semantics beyond acquire and releases, this is sequentially consistent
  - Consumed=dependency-based ordering



# Std::atomic\_bool

- Most basic `std::atomic_bool`, can be built from a non-atomic bool
- Can be constructed, initialized, assigned from a plain bool
- **assignment operator from a non-atomic bool does not return a reference to the object assigned to, but it returns a bool with the value assigned (like all other atomic types).**
  - prevents code that depended on the result of the assignment to have to explicitly load the value, potentially getting a modified result from another thread.
- **replace the stored value with a new one and retrieve the original value**
- **a plain non-modifying query of the value with an implicit conversion to plain bool**
- **RMW operation that stores a new value if the current value is equal to an expected value is `compare_exchange_{weak/strong}()`;**
- **If we have spurious failure:**

```
bool expected=false; extern atomic_bool b; // set somewhere else
while(!b.compare_exchange_weak(expected,true) && !expected);
```

- **May not be lock free, need to check per instance**

## Std::atomic\_address

- **similar operations as std::atomic\_bool, can be constructed from non-atomic void\***
- **all operations take and return void\***
- **adds operations fetch\_add() and fetch\_sub() and += and -= as wrappers**

# Basic atomics

- **atomic<bool>**
  - Load, store, swap, cas
- **atomic<int>**
  - Load, store, swap, cas
  - Fetch-and-(add, sub, and, or, xor)
- **atomic<void \*>**
  - Load, store, swap, cas
  - Fetch-and-(add, sub)

# Std::atomic\_integral

- this adds `fetch_and`, `fetch_or`, `fetch_xor`, and compound assignments like:
- `+=`, `-=`, `&=`, `^=`, pre and post increment and decrement
- missing division, multiplication and shift operations, but atomic integrals are usually used as counters or bit masks, this is not a big loss
- all semantics match `fetch_add` and `fetch_sub` for `atomic_address`: returns old value
- the compound assignments return new value
- `++x` increments the variable and returns new value, `x++` increments the variable and returns old value
- result is the value of the associated integral type

## Other Atomic facilities

- Specializations for integral types, pointers
  - **Provide atomic increment, decrement (++ , -- , += , -=)**
  - **Note: `x++` is very different from `x = x + 1` !**
    - **Unlike Java volatiles, where both are probably wrong!**
- Non-template (C-like) atomic types
  - **Template specializations inherit from these**
- C-like stand-alone (`atomic_`) function interfaces.

# Std::atomic <> template

- **std::atomic<>** to create an atomic user-defined type
- Specializations for integral types derived from `std::atomic_integral_type`, and pointer types
- Main benefit of the template is atomic variants of user-defined types, can't be just any UDT, it must fit this criteria:
  - must have trivial copy-assignment operator: no virtual functions or virtual bases and must use the compiler-generated copy-assignment operator
  - every base class and non-static data member of UDT must also have a trivial copy-assignment operator
  - Must be bitwise equality comparable
- **Only have**
  - `load()`, `store()`
- **Assignment and conversion to the UDT**
  - `exchange()`, `compare_exchange_weak()`, `compare_exchange_strong()`
  - assignment from and conversion to an instance of type T

# Free functions

- **Designed to be C compatible, so they use pointers and not references**
- **overloaded for each atomic type**
- **all take a pointer to the atomic object as first parameter**
- **2 varieties**
  - one without the memory order tag
  - one with an `_explicit` suffix and additional memory ordering tag, or tags
- **`std::atomic_is_lock_free()` comes in only one variety**
  - `std::atomic_is_lock_free(&a)` returns the same value as `a.is_lock_free()`
- **`std::atomic_load(&a)` returns the same value as `a.load()`;**
- **`std::atomic_load_explicit(&a, std::memory_order_acquire)` is the same as `a.load(std::memory_order_acquire)`**



# Sequencing redefined for serial program

- **Sequence points are ... gone!**
- **Sequence are now defined by ordering relations**
  - Sequence-before
  - Indeterminately-sequenced
- **A write/write or read/write pair relations**
  - That are not sequenced before
  - That are not indeterminately-sequenced
  - Results in undefined behaviour

# Sequencing extended for parallel programs

- **Sequenced-before**
  - Provides intra-thread ordering
- **Synchronizes with (Acquire and release)**
  - Provide inter-thread ordering
- **Happens-before relation**
  - Between memory operations in different threads

# Sequenced before

- **If a memory update or side-effect *a* is-sequenced-before another memory operation or side-effect *b*,**
  - then informally *a* must appear to be completely evaluated before *b* in the sequential execution of a single thread, e.g. all accesses and side effects of *a* must occur before those of *b*.
  - We will say that a subexpression *A* of the source program *is-sequenced-before* another subexpression *B* of the same source program to indicate that all side-effects and memory operations performed by an execution of *A* occur-before those performed by the corresponding execution of *B*, i.e. as part of the same execution of the smallest expression that includes them both.
- **We propose roughly that wherever the current standard states that there is a sequence point between *A* and *B*, we instead state that *A* is-sequenced-before *B*. This will constitute the precise definition of *is-sequenced-before* on subexpressions, and hence on memory actions and side effects.**

# Cases

- **Function calls:**

- The evaluations of the postfix expression and of the argument expressions are all unsequenced relative to one another. All side effects of argument expression evaluations are sequenced before the function is entered

- **Increment & Decrement:**

- The value computation of the ++ expression is sequenced before the modification of the operand object.

- **Logical AND operator**

- If the second expression is evaluated, every value computation and side effect associated with the first expression is sequenced before every value computation and side effect associated with the second expression.

- **Conditional Operator**

- Every value computation and side effect associated with the first expression is sequenced before every value computation and side effect associated with the second or third expression.

- **Comma operator**

- Every value computation and side effect associated with the left expression is sequenced before every value computation and side effect associated with the right expression.

# Synchronizes with

- only between operations on atomic types
- operations on a data structure ( locking a mutex) might provide this relationship if the data structure contains atomic types, and the operations on that data structure perform the appropriate operations internally
- definition:
  - **a suitably-tagged atomic write operation on a variable x synchronizes-with a suitably-tagged atomic read operation on x that reads the value stored**  
by (a) that write, (b) a subsequent atomic write operation on x by the same thread that performed the initial write, or (c) an atomic read-modify-write operation on x (such as `fetch_add()` or `compare_exchange_weak()`) by any thread, that read the value written.
- Store-release synchronizes-with a load-acquire

# Happens before

- It specifies which operations see the effects of which other operations.
- An evaluation A happens before an evaluation B if:
  - **A is sequenced before B, or**
  - **A synchronizes with B, or**
  - **for some evaluation X, A happens before X and X happens before B.**

# Happens-before

*Thread 1*

```

if (!x_init.ld_acq())
{
    lock();
    if (!x_init.ld_...())
        x = ...;
    x_init.store_rel(1);
    unlock();
}
... x ...

```

*Thread 2*

```

if (!x_init.ld_acq())
{
    lock();
    if (!x_init.ld_...())
        x = ...;
    x_init.store_rel(1);
    unlock();
}
... x ...

```

#3

#4

#1

#2



# Memory Ordering Operations

```
typedef enum memory_order {  
    memory_order_relaxed, memory_order_consume,  
    memory_order_acquire,  
  
    memory_order_release, memory_order_acq_rel,  
    memory_order_seq_cst  
} memory_order;
```

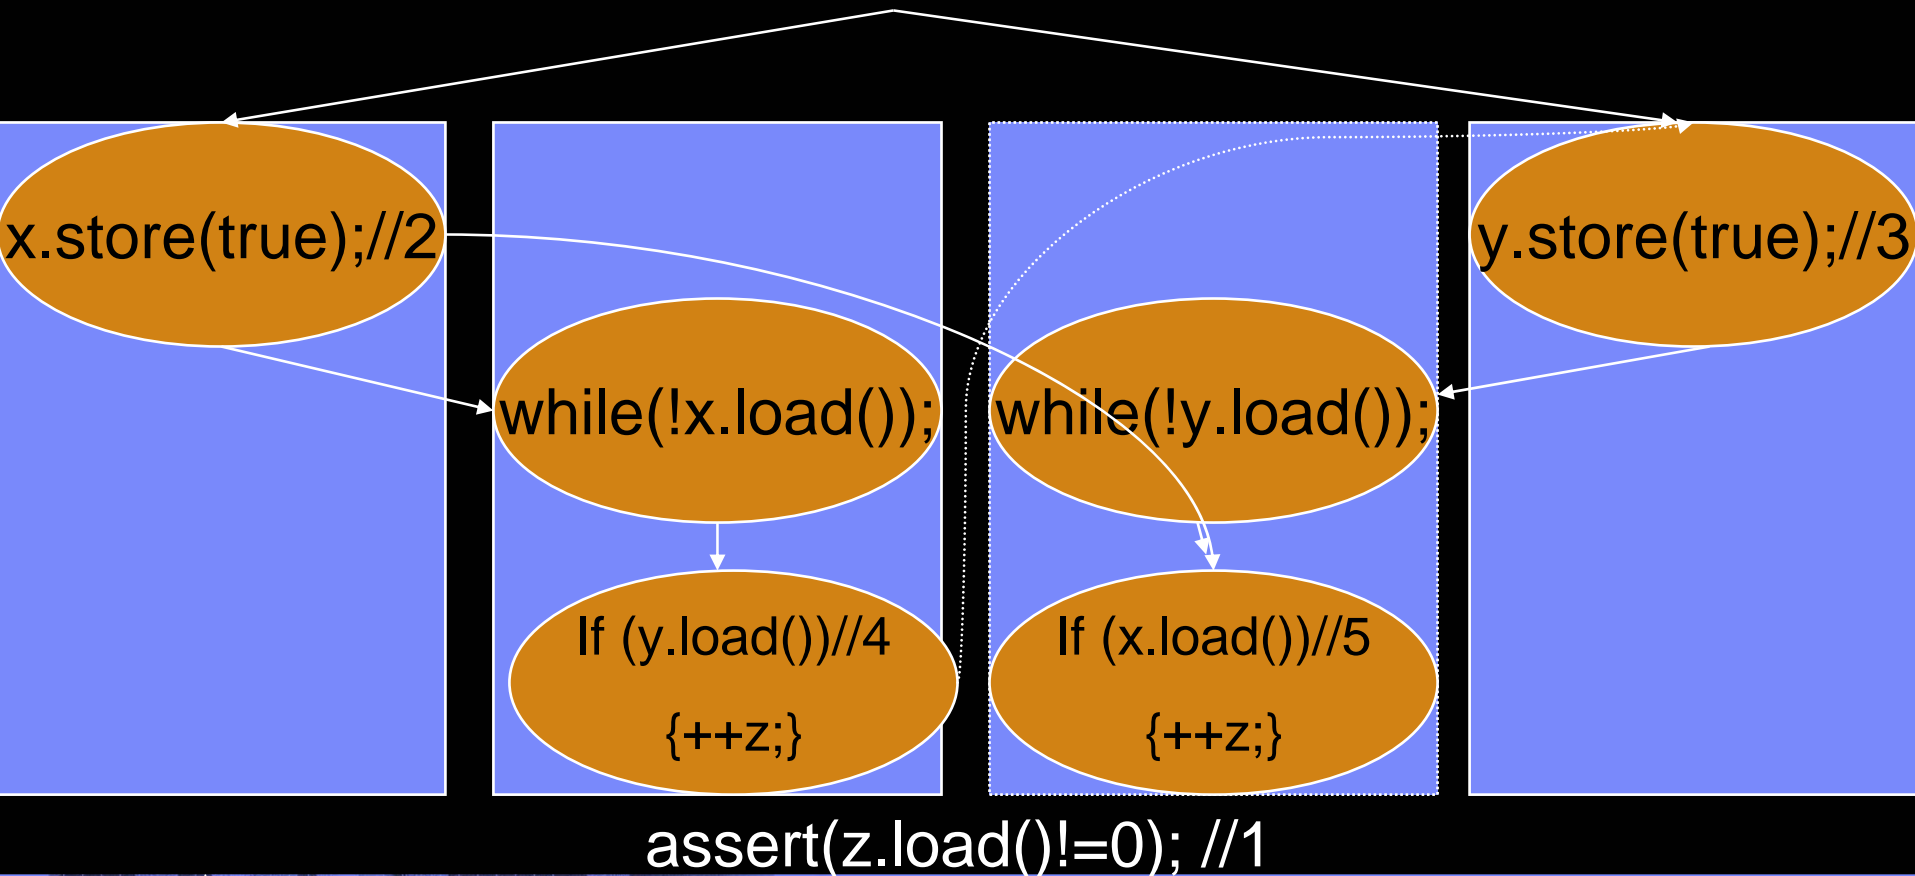
- **Every atomic operation has a default form, implicitly using `seq_cst`, and a form with an explicit order argument**
- **When specified, argument is expected to be just an enum constant**

# Memory Ordering Constraints

- Sequential Consistency
  - **Single total order for all SC ops on all variables**
  - **default**
- Acquire/Release
  - **Pairwise ordering rather than total order**
  - **Independent Reads of Independent Writes don't require synchronization between CPUs**
- Relaxed Atomics
  - **Read or write data without ordering**
  - **Still obeys happens-before**

# SC and happens-before

```
std::atomic_bool x,y; std::atomic_int z;x=false;y=false;z=0;
```



# Relaxed and happens-before

```
std::atomic_bool x,y; std::atomic_int z;x=false;y=false;z=0;
```

x.store(true); //4

y.store(true); //5

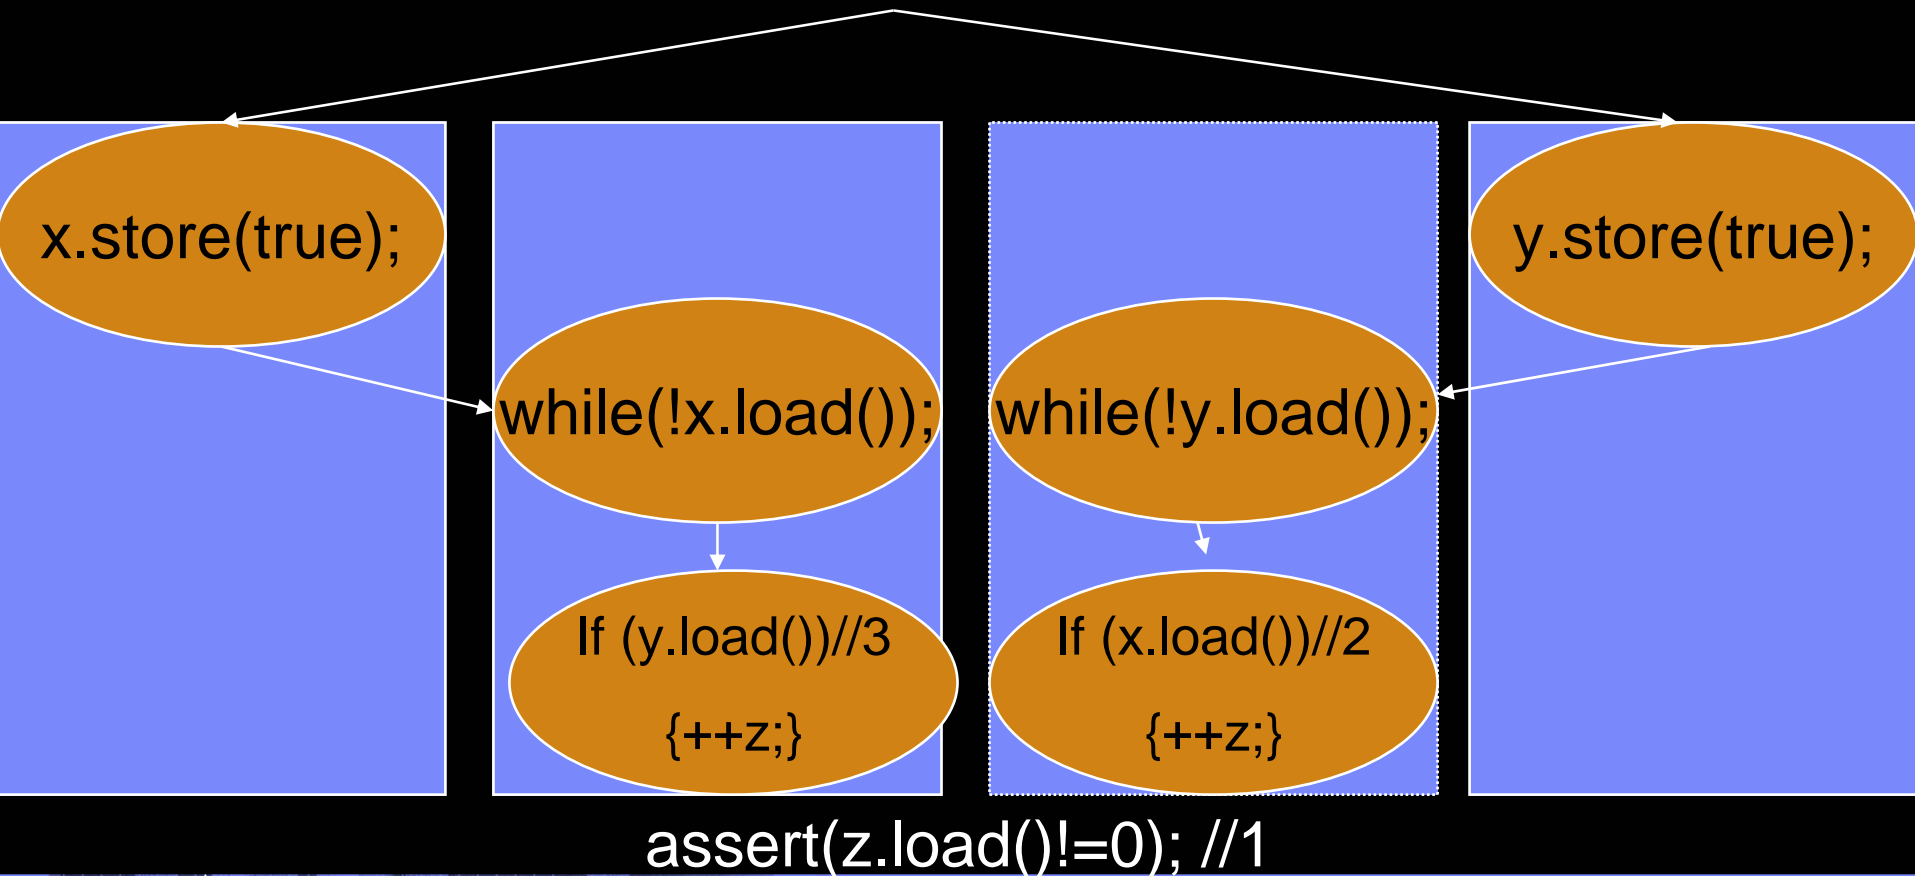
while(!y.load()); //3

If (x.load()) //2  
{ ++z; }

assert(z.load() != 0); //1

# Acquire-Release with Happens-before

```
std::atomic_bool x,y; std::atomic_int z;x=false;y=false;z=0;
```



# Food for thought and Q/A

- **This is the chance to make comments on the C++0x FCD draft through us or the National Body rep:**
  - <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2010/n3092.pdf>
- **Memory Model:**
  - [http://www.hpl.hp.com/personal/Hans\\_Boehm/c++mm](http://www.hpl.hp.com/personal/Hans_Boehm/c++mm)
- **Participate and feedback to Compiler**
- **Talk to me at my blog:**
  - <http://www.ibm.com/software/rational/cafe/blogs/cpp-standard>

# My blogs and email address

- michaelw@ca.ibm.com
- Rational C/C++ cafe: <http://www.ibm.com/software/rational/cafe/community/ccpp>
- My Blogs:
- Parallel & Multi-Core Computing multicore <http://www.ibm.com/software/rational/cafe/blogs/ccpp-parallel->
- C++ Language & Standard <http://www.ibm.com/software/rational/cafe/blogs/cpp-standard>
- Commercial Computing commercial <http://www.ibm.com/software/rational/cafe/blogs/ccpp->
- Boost test results  
<http://www.ibm.com/support/docview.wss?rs=2239&context=SSJT9L&uid=swg27006911>
- C/C++ Compilers Support Page <http://www.ibm.com/software/awdtools/ccompilers/support/>
- C/C++ Feature Request Interface <http://www.ibm.com/support/docview.wss?uid=swg27005811>
- XL Fortran Compiler Support Page  
<http://www.ibm.com/software/awdtools/fortran/xlfortran/support/>
- XL Fortran Feature Request Interface <http://www.ibm.com/support/docview.wss?uid=swg27005812>

# Acknowledgement

- **Some slides are borrowed from committee presentations by various committee members, their proposals, and private communication**