

# Transactional Memory at BoostCon'10



*BoostCon 2010*

*Vicente Botet, Justin E. Gottschlich, Maurice Herlihy,  
Mark Moir, Tatiana Shpeisman, Michael Wong*

# TM Lineup



Vicente Botet



Maurice Herlihy



Justin Gottschlich




Tatiana Shpeisman

Michael Wong



Mark Moir



Tatiana (Part II)



Today

- Library
  - TBoost.STM Engine
- Compiler
  - Proposed C++ language extension
  - IBM
  - Sun / Oracle
  - Intel



- Experts panel
  - Maurice Herlihy
  - Mark Moir
  - Tatiana Shpeisman
  - Michael Wong
- Submit questions today



# *Why Should You Care?*

- May affect you
  - C++ language extension
  - Library / Compiler
- Insight
  - Understand parallel algorithms
  - Understand how TM applies
- Contacts

# The TBoost.STM Engine



## *BoostCon 2010*

*Justin E. Gottschlich*

University of Colorado-Boulder

**Raytheon**

# Motivation

- Problem
  - TM is not fast enough! (Cascaval et al., 2008)
- Reason
  - Conflict Detection and Opacity
  - Most TMs use *Validation*
- Our solution:
  - *Full Invalidation*
  - *InvalSTM*





# TM Performance Bottleneck



- *Conflict Detection*

- Determine if transaction can commit
  - (Papadimitrou, "Theory of Database Concurrency Control," 1986)

- *Opacity*

- Keep in-flight transactions consistent
  - (Guerraoui & Kapalka, PPOPP'08)



# Conflict Detection

*Conflict:*  $W_{T1} \cap (W_{T2} \cup R_{T2}) \neq \emptyset$



- **Validation (T2)**

- *Analyze the Past*

- Version # is same

- **Invalidation (T1)**

- *Analyze the Future*

- $T2.valid = false$

# Opacity



- **Validation**

- Version # is same

- **Invalidation**

- Check *valid*  $\neq$  *false*

# Opacity: Validation

```
atomic { return (x == 0 ? 0 : y / z); }
```

read x → store x's ver / ref

read y → store y's ver / ref,  
*opacity*: validate x

read z → store z's ver / ref,  
*opacity*: validate x, y

commit → *conflict*: validate x, y, z

time

# Opacity: Invalidation

```
atomic { return (x == 0 ? 0 : y / z); }
```

read x → store ref to x

read y → store ref to y

*opacity*: if (!*valid*) retry

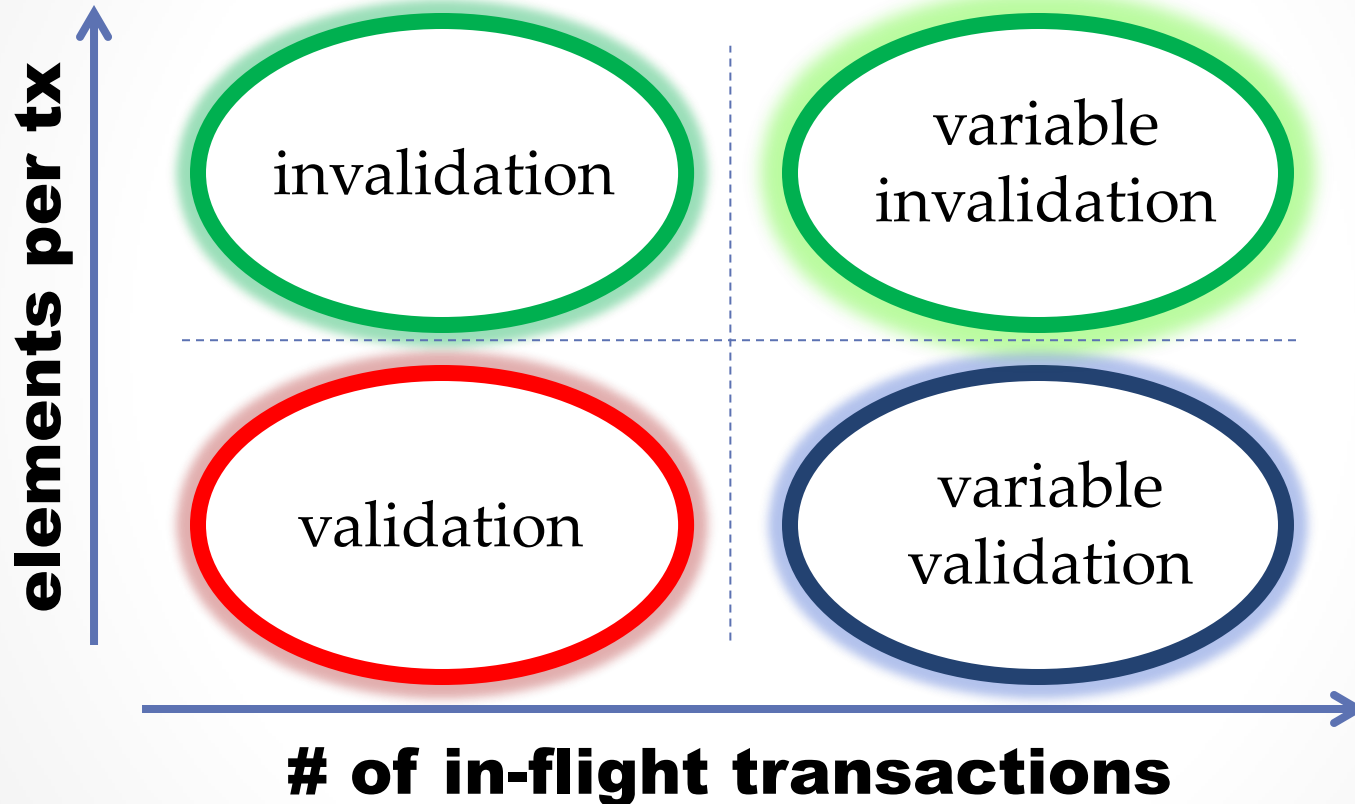
read z → store ref to z

*opacity*: if (!*valid*) retry

commit → *conflict*: none


time

# Validation Vs. Invalidation



# Contending + Concurrent Workload

## 1-Writer, N-Reader




Commit to Executed Ratio: *Commits / Executed*

Max = 1, Min = 0

# Side-By-Side Analysis


*Validation*



Commit / Executed:  $1 / M$

$$\lim_{M \rightarrow \infty} \left( \frac{1}{M} \right) = 0$$

*Invalidation*



Commit / Executed:  $(M-1) / M$

$$\lim_{M \rightarrow \infty} \left( \frac{(M-1)}{M} \right) = 1$$



# Algorithmic Growth

$$\textit{Validation} = \sum_{i=1}^M \sum_{j=1}^{r_i} j$$

$$\textit{Invalidation} = \sum_{i=1}^M \left( r_i + \sum_{j=1}^{F_i} w_i (s_{rj}(r_j) + s_{wj}(w_j)) \right)$$

$$\textit{Bloom Inval} = \sum_{i=1}^M (r_i + (2kw * Fi))$$

# Efficient Read-Only Transactions

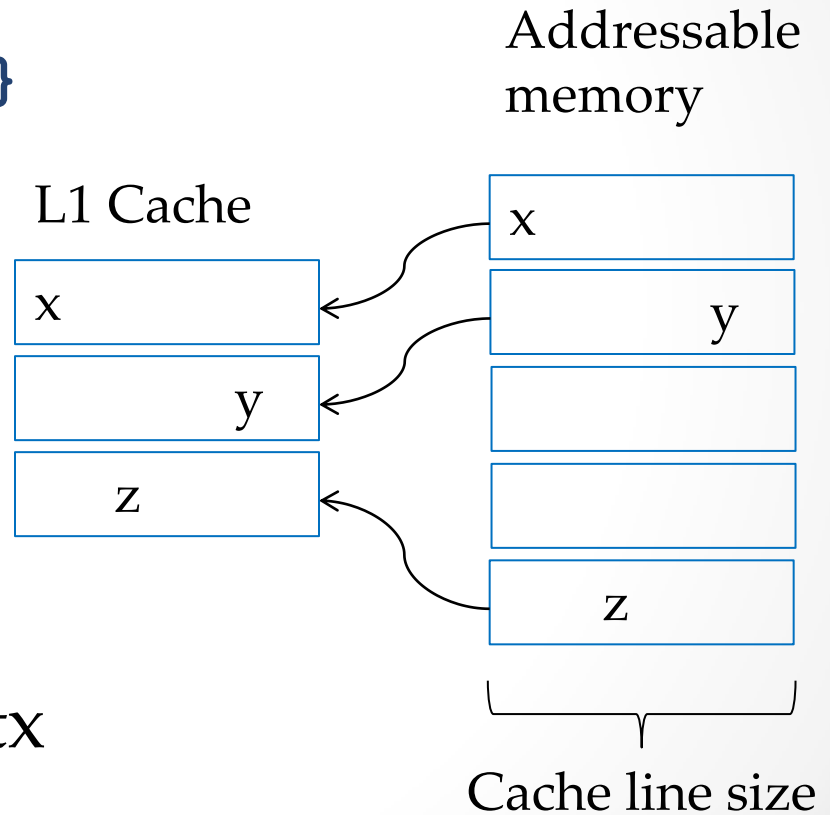
$$\textit{Validation Read-Only} = \sum_{i=1}^M \sum_{j=1}^{r_i} j$$

$$\textit{Invalidation} = \sum_{i=1}^M \left( r_i + \sum_{j=1}^{F_i} w_i(s_{r_j}(r_j) + s_{w_j}(w_j)) \right)$$

$$\textit{Invalidation Read-Only} = \sum_{i=1}^M r_i$$

# Validation + Memory

```
atomic { x = y / z; }
```

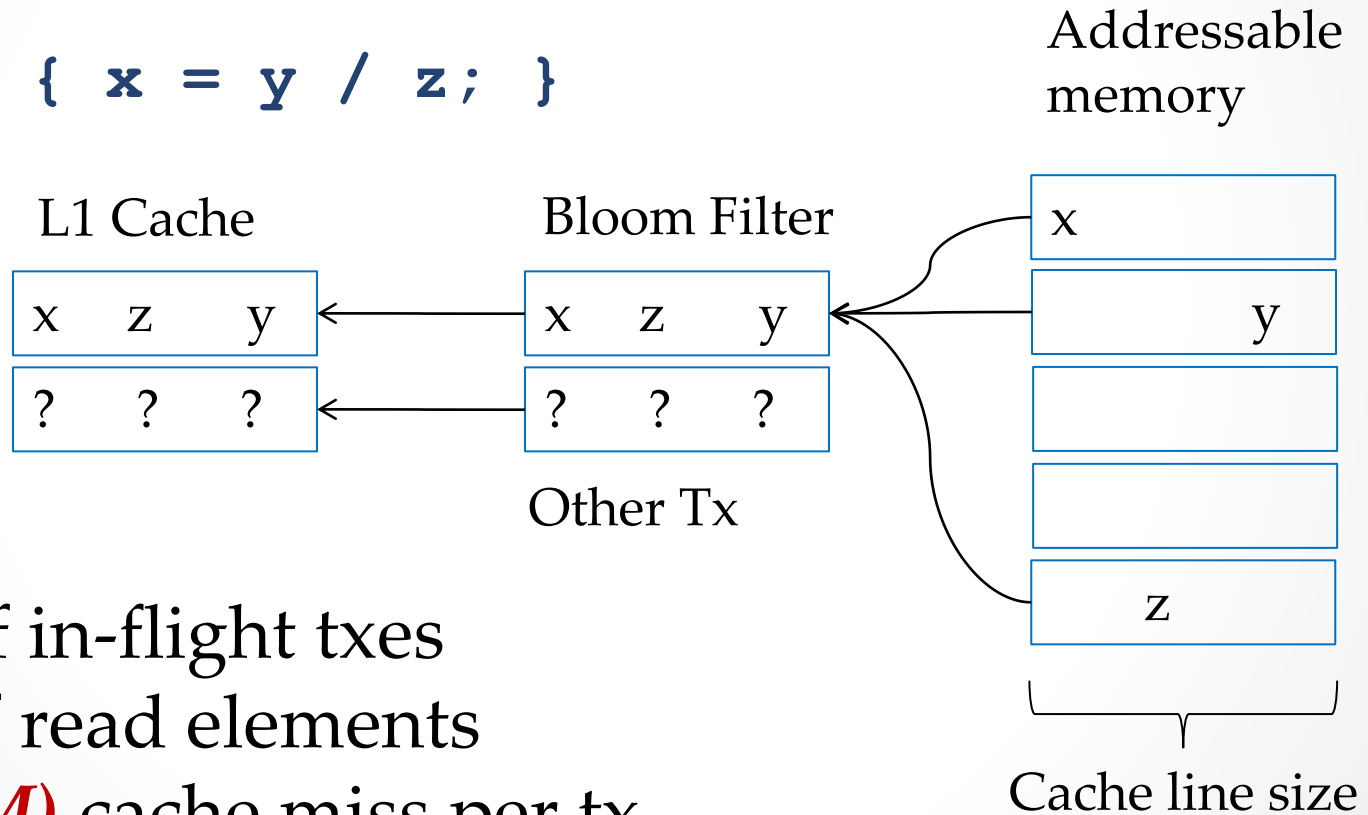


**N** = elements per tx

**$O(N^2)$**  cache misses per tx

# Invalidation + Memory

```
atomic { x = y / z; }
```




$M$  = # of in-flight txes


$N$  = # of read elements

$O(N + M)$  cache miss per tx


# Linked List



# 1-Writer / N-Readers




# Hash Table



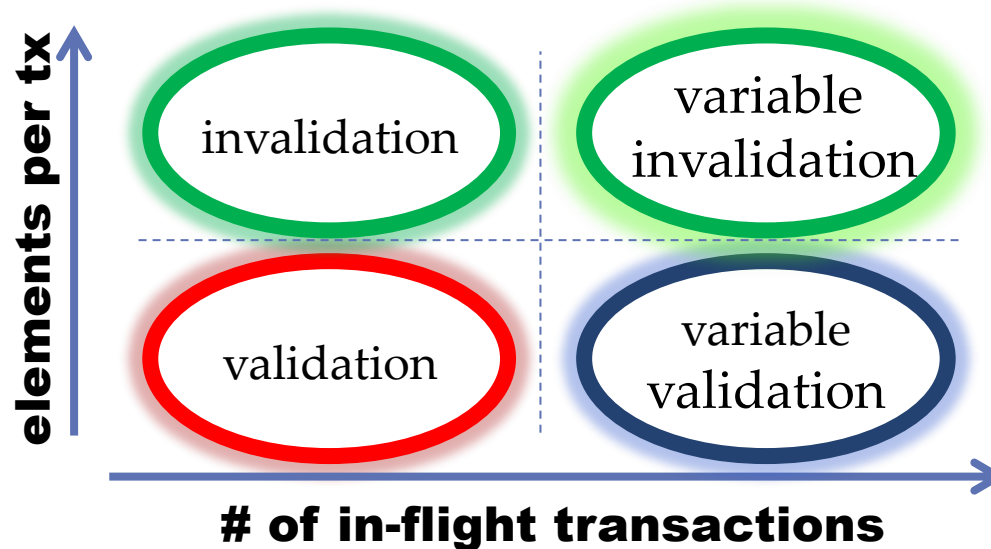


# Zoomed Hash Table



# Conclusion

- InvalSTM (TBoost.STM) competitive with TL2
- Invalidation (InvalSTM) can be efficient



- Next up
  - Proof of correctness for Full Invalidation
  - InvalSTM + STAMP
- Special thanks to Spear and Herlihy

# Questions?



Justin E. Gottschlich

[gottschl@colorado.edu](mailto:gottschl@colorado.edu)

<http://eces.colorado.edu/~gottschl/>

# Panel Questions

What is the most important  
TM problem to solve in the  
next 1-2 years?

# Panel Questions

What is the most important  
TM problem to solve in the  
next 5-10 years?

# Panel Questions

Is TM on track to becoming a real parallel programming solution?

If so, when will it be realized?

If not, what are we doing wrong and what do we need to do to fix it?

# Panel Questions

In what types of environments (systems or programs) do you see TM being most useful?

Why?



# Panel Questions

Where is TM least useful?

Why?

# Panel Questions

What will parallel programs  
look like in the future?

# Panel Questions

What should a transaction do  
when an exception occurs  
within it?

Abort? Commit?